

Part 1

Soln1)

```
From sklearn.datasets import fetch_california_housing
```

```
From sklearn.linear_model import LinearRegression
```

```
Import matplotlib.pyplot as plt
```

```
# Load dataset
```

```
Housing = fetch_california_housing()
```

```
Rooms = housing.data[:, 3].reshape(-1, 1) # AveRooms
```

```
Prices = housing.target # Median House Value
```

```
# Train linear regression model
```

```
Model = LinearRegression()
```

```
Model.fit(rooms, prices)
```

```
# Predict and plot
```

```
Predicted_prices = model.predict(rooms)
```

```
Plt.scatter(rooms, prices, color='skyblue', alpha=0.5, label='Actual Prices')
```

```
Plt.plot(rooms, predicted_prices, color='crimson', linewidth=2, label='Regression Line')
```

```
Plt.xlabel('Average Rooms per Household')
```

```
Plt.ylabel('Median House Value')
```

```
Plt.title('Simple Linear Regression')
```

```
Plt.legend()
```

```
Plt.show()
```

Soln2)

```
From sklearn.model_selection import train_test_split

From sklearn.metrics import r2_score, mean_squared_error

Import numpy as np


# Select numeric features: MedInc, HouseAge, AveRooms, AveOccup
Features = housing.data[:, [0, 1, 3, 5]]

Target = housing.target

# Split dataset
X_train, X_test, y_train, y_test = train_test_split(features, target, test_size=0.2,
random_state=42)


# Train model
Model = LinearRegression()
Model.fit(X_train, y_train)


# Predict and evaluate
Y_pred = model.predict(X_test)

R2 = r2_score(y_test, y_pred)

Mse = mean_squared_error(y_test, y_pred)

Rmse = np.sqrt(mse)

Print("R-squared:", r2)

Print("Mean Squared Error:", mse)

Print("Root Mean Squared Error:", rmse)

Print("Feature Coefficients:", model.coef_)
```

Soln 3)

```

from sklearn.preprocessing import StandardScaler

model_unscaled = LinearRegression()

model_unscaled.fit(X_train, y_train)

r2_unscaled = model_unscaled.score(X_test, y_test)

# With scaling

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)

X_test_scaled = scaler.transform(X_test)

model_scaled = LinearRegression()

model_scaled.fit(X_train_scaled, y_train)

r2_scaled = model_scaled.score(X_test_scaled, y_test)

print("R2 without Scaling:", r2_unscaled)

print("R2 with Scaling:", r2_scaled)

```

soln 4)

```

import pandas as pd

import seaborn as sns

import matplotlib.pyplot as plt

# Convert data to DataFrame for easier analysis

Df = pd.DataFrame(housing.data, columns=housing.feature_names)

Df['MedHouseVal'] = housing.target # Add target column

# Compute correlation matrix

Corr_matrix = df.corr()

# Plot heatmap

```

```
Plt.figure(figsize=(10, 8))  
  
Sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f")  
  
Plt.title('Correlation Matrix Heatmap')  
  
Plt.show()
```

Part 2

Soln 5)

```
From sklearn.datasets import load_breast_cancer  
  
From sklearn.linear_model import LogisticRegression  
  
From sklearn.model_selection import train_test_split  
  
From sklearn.metrics import accuracy_score, confusion_matrix, classification_report,  
roc_auc_score, roc_curve  
  
Import matplotlib.pyplot as plt  
  
  
# Load dataset  
  
Data = load_breast_cancer()  
  
X = data.data  
  
Y = data.target  
  
  
# Split dataset  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)  
  
  
# Train logistic regression model  
  
Model = LogisticRegression(max_iter=1000)  
  
Model.fit(X_train, y_train)
```

```

# Predict and evaluate

Y_pred = model.predict(X_test)

Y_probs = model.predict_proba(X_test)[:, 1]


# Evaluation metrics

Print("Accuracy:", accuracy_score(y_test, y_pred))

Print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

Print("Classification Report:\n", classification_report(y_test, y_pred))

Print("ROC-AUC Score:", roc_auc_score(y_test, y_probs))


# Plot ROC Curve

Fpr, tpr, thresholds = roc_curve(y_test, y_probs)

Plt.plot(fpr, tpr, label="ROC Curve")

Plt.plot([0, 1], [0, 1], 'k—') # Diagonal

Plt.xlabel("False Positive Rate")

Plt.ylabel("True Positive Rate")

Plt.title("ROC Curve")

Plt.legend()

Plt.show()

```

Soln 6)

```

From sklearn.metrics import f1_score

```

For threshold in [0.3, 0.5, 0.7]:

```

Custom_pred = (y_probs >= threshold).astype(int)

```

```

Cm = confusion_matrix(y_test, custom_pred)

```

```

F1 = f1_score(y_test, custom_pred)

Print(f"Threshold: {threshold}")

Print("Confusion Matrix:\n", cm)

Print("F1 Score:", f1)

Print("-" * 30)

# Plot ROC Curve with threshold marker

Plt.plot(fpr, tpr, label="ROC Curve")

Plt.scatter(fpr[np.argmax(tpr - fpr)], tpr[np.argmax(tpr - fpr)],
            Color='red', label='Optimal Threshold')

Plt.xlabel("False Positive Rate")

Plt.ylabel("True Positive Rate")

Plt.title("ROC Curve with Optimal Threshold")

Plt.legend()

Plt.show()

```

PART 3

1) Linear Regression Assumptions: Linearity, independence of errors, homoscedasticity, and normally distributed residuals.

2) Use Logistic Regression: When the target variable is categorical, especially binary (e.g., yes/no).

3) Coefficients in Logistic Regression: Indicate the change in log odds for a one-unit increase in the predictor.

4) Sigmoid vs. Softmax: Sigmoid outputs a single probability (binary); softmax gives a probability distribution across multiple classes.

5) R-squared in Logistic Models: It doesn't apply because logistic regression predicts probabilities, not continuous values.